

Project Milestone 3 - Data Storage Implementation: KV + relational

Esam Uddin - 100711116, Mihir Patel - 100702168, Jane Coralde 100660214, Haider Sarmad

3/18/2022

Given Lab 3 Repository: <https://github.com/goergeddaoud/SOFE4630U-tut3>

Group Project Repository:

<https://github.com/esam191/Intelligent-Transportation-System>

Objectives:

- Deploy Tabular and key-Value data storage to GKE.
- Get familiar with Key-Value data storage
- Get familiar with Kafka Connectors and their configuration.
- Configure and use Kafka source connector to Redis.
- Configure and use MySQL sink and source Kafka connectors.

Procedure:

- 1. Watch the first three videos for Kafka connectors (focus on the concepts, not the details).**
- 2. Describe the following:**
 - **Sink and Source connectors.**
 - The source connector basically imports data from any relational database like MySQL with another driver into a Kafka topic.
 - Source connector also collects and stores metrics from application servers in Kafka topics.
 - On the other hand, the sink connector exports data from Kafka topics to any relational database with the help of a driver/secondary indexes like Elasticsearch.
 - In other words, both connectors are essentially just used to copy data between Kafka and other systems where you want to push

data to or pull data from.

- **The applications/advantages of using Kafka Connectors with data storage.**
 - Using Kafka connectors allows you to stream data both scalably and reliably between Kafka and other data storage systems.
 - Kafka connectors help with moving large data sets both into and out of Kafka.
 - Kafka connectors also allow for a data centric pipeline in which appropriate data abstractions are used.
 - Using Kafka connectors also provides flexibility as data streaming runs on a single node.
- **How do Kafka connectors maintain availability?**
 - Kafka connectors maintain availability by running the streaming and batch-oriented systems on a single node.
 - Kafka connectors provide distributed service, which basically means it is scaled to an organization wide service.
 - Kafka connectors provide lower time to production by using existing connectors (avoids downtime), which in turn provides availability.
- **List the popular Kafka converters for values and the properties/advantages of each.**
 - JsonConverter
 - Embeds schema in the JSON
 - Serializes message keys and values into JSON docs
 - AvroConverter
 - Specifies schema registry
 - Can be used for message payload
 - ProtobufConverter
 - performs protobuf serialization
 - StringConverter
 - Just a string
 - No schema to the data

3. Search the internet to answer the following question:

- **What's a Key-Value (KV) database?**
 - Provides storing, retrieving of key value data
 - Known as a dictionary or hash table
 - Records are stored and retrieved using a unique key

Ex. table shows values associated with keys

Key	Value
K1	AA,FD,FG
K2	FS,GG,HH
K3	1234,SAS
K4	FF,12,2222

- **What are KV databases' advantages and disadvantages?**
 - Advantages
 - Scalability: scalable both vertically and horizontally
 - Reliability: built-in redundancy
 - Speed: quick to respond, data has any type/multiple types
 - Simplicity: simple to use
 - Disadvantages
 - No query language: difficult to transport queries from one database to a KV database
 - No filtration of values: whole values are returned when a request is made
 - Too simple: not very refined
- **List some popular KV databases.**
 - Redis
 - Amazon DynamoDB
 - Oracle NoSQL DB
 - InfinityDB
 - Aerospike

4. Follow the following videos to deploy and use Redis and MySQL databases using GKE.

Video 1: <https://www.youtube.com/watch?v=qVD1uVKMZYc>

```
esaml91@cloudshell:~/SOFE4630U-tut3/GKE (silver-course-344506) $ kubectl apply -f mysql-pvc.yaml
persistentvolumeclaim/mysql-volumeclaim created
esaml91@cloudshell:~/SOFE4630U-tut3/GKE (silver-course-344506) $ kubectl get pvc
NAME                                STATUS    VOLUME                                     CAPACITY   ACCESS MODES   STORAGECLASS   AGE
mysql-volumeclaim                  Bound    pvc-99c2054b-145d-4368-bf17-494489544a17  10Gi       RWO             standard       10s
esaml91@cloudshell:~/SOFE4630U-tut3/GKE (silver-course-344506) $ kubectl apply -f mysql-app.yaml
service/mysql created
deployment.apps/mysql created
esaml91@cloudshell:~/SOFE4630U-tut3/GKE (silver-course-344506) $ kubectl get pods
NAME                                READY    STATUS    RESTARTS   AGE
mysql-7dcb5fd764-zjj6s              0/1      ContainerCreating    0          18s
esaml91@cloudshell:~/SOFE4630U-tut3/GKE (silver-course-344506) $ kubectl get pods
NAME                                READY    STATUS    RESTARTS   AGE
mysql-7dcb5fd764-zjj6s              1/1      Running    0          34s
esaml91@cloudshell:~/SOFE4630U-tut3/GKE (silver-course-344506) $ kubectl get deployment
NAME    READY   UP-TO-DATE   AVAILABLE   AGE
mysql  1/1     1            1           54s
esaml91@cloudshell:~/SOFE4630U-tut3/GKE (silver-course-344506) $ kubectl get services
NAME    TYPE        CLUSTER-IP    EXTERNAL-IP    PORT(S)    AGE
kubernetes  ClusterIP   10.12.0.1     <none>         443/TCP    4m34s
mysql     LoadBalancer 10.12.14.37   104.197.66.164 3306:30768/TCP 74s
esaml91@cloudshell:~/SOFE4630U-tut3/GKE (silver-course-344506) $
```

Video 2: <https://www.youtube.com/watch?v=GBvGjVLbYIs>

```
mysql> exit
Bye
esaml91@cloudshell:~/SOFE4630U-tut3/GKE (silver-course-344506) $ ls
mysql-app.yaml mysql-pvc.yaml redis-app.yaml redis-pvc.yaml scl.sql
esaml91@cloudshell:~/SOFE4630U-tut3/GKE (silver-course-344506) $ cat scl.sql
CREATE DATABASE IF NOT EXISTS myDB;
USE myDB;

DROP TABLE IF EXISTS test;

CREATE TABLE IF NOT EXISTS test (
  id serial NOT NULL PRIMARY KEY,
  name varchar(100),
  email varchar(200),
  department varchar(200),
  modified timestamp default CURRENT_TIMESTAMP NOT NULL,
  INDEX `modified_index` (`modified`)
);
USE myDB;
INSERT INTO test (name, email, department) VALUES ('alice', 'alice@abc.com', 'eng.');
```

```

esam191@cloudshell:~/SOFE4630U-tut3/GKE (silver-course-344506)$ cat sc1.sql | kubectl exec -i mysql-7dcb5fd764-zjj6s -- mysql -uuser -pSOFE4630U
mysql: [Warning] Using a password on the command line interface can be insecure.
esam191@cloudshell:~/SOFE4630U-tut3/GKE (silver-course-344506)$ echo "select * from myDB.test;" | kubectl exec -i mysql-7dcb5fd764-zjj6s -- mysql -uuser -pSOFE4630U
-bash: kubectl: command not found
esam191@cloudshell:~/SOFE4630U-tut3/GKE (silver-course-344506)$ echo "select * from myDB.test;" | kubectl exec -i mysql-7dcb5fd764-zjj6s -- mysql -uuser -pSOFE4630U
mysql: [Warning] Using a password on the command line interface can be insecure.

```

id	name	email	department	modified
1	alice	alice@abc.com	eng.	2022-03-18 07:18:09
2	bob1	bob1@abc.com	sales	2022-03-18 07:18:09
3	bob2	bob2@abc.com	sales	2022-03-18 07:18:09
4	bob3	bob3@abc.com	sales	2022-03-18 07:18:09
5	bob4	bob4@abc.com	sales	2022-03-18 07:18:09
6	bob5	bob5@abc.com	sales	2022-03-18 07:18:09
7	bob6	bob6@abc.com	sales	2022-03-18 07:18:09
8	bob7	bob7@abc.com	sales	2022-03-18 07:18:09
9	bob8	bob8@abc.com	sales	2022-03-18 07:18:09
10	bob9	bob9@abc.com	sales	2022-03-18 07:18:09

```

mysql> use myDB;
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A

Database changed
mysql> select * from test;

```

id	name	email	department	modified
1	alice	alice@abc.com	eng.	2022-03-18 07:18:09
2	bob1	bob1@abc.com	sales	2022-03-18 07:18:09
3	bob2	bob2@abc.com	sales	2022-03-18 07:18:09
4	bob3	bob3@abc.com	sales	2022-03-18 07:18:09
5	bob4	bob4@abc.com	sales	2022-03-18 07:18:09
6	bob5	bob5@abc.com	sales	2022-03-18 07:18:09
7	bob6	bob6@abc.com	sales	2022-03-18 07:18:09
8	bob7	bob7@abc.com	sales	2022-03-18 07:18:09
9	bob8	bob8@abc.com	sales	2022-03-18 07:18:09
10	bob9	bob9@abc.com	sales	2022-03-18 07:18:09

```

10 rows in set (0.00 sec)

```

```

esam191@cloudshell:~/SOFE4630U-tut3/GKE (silver-course-344506)$ mysql -uuser -pSOFE4630U -h34.134.216.154 <<< "select * from myDB.test"
mysql: [Warning] Using a password on the command line interface can be insecure.

```

id	name	email	department	modified
1	alice	alice@abc.com	eng.	2022-03-18 07:28:42
2	bob1	bob1@abc.com	sales	2022-03-18 07:28:42
3	bob2	bob2@abc.com	sales	2022-03-18 07:28:42
4	bob3	bob3@abc.com	sales	2022-03-18 07:28:42
5	bob4	bob4@abc.com	sales	2022-03-18 07:28:43
6	bob5	bob5@abc.com	sales	2022-03-18 07:28:43
7	bob6	bob6@abc.com	sales	2022-03-18 07:28:43
8	bob7	bob7@abc.com	sales	2022-03-18 07:28:43
9	bob8	bob8@abc.com	sales	2022-03-18 07:28:43
10	bob9	bob9@abc.com	sales	2022-03-18 07:28:43

```

esam191@cloudshell:~/SOFE4630U-tut3/GKE (silver-course-344506)$

```

Video 3: <https://www.youtube.com/watch?v=9R3eEAlpOtk&feature=youtu.be>

```
esam191@cloudshell:~/SOFE4630U-tut3/GKE (silver-course-344506)$ kubectl get services
NAME          TYPE          CLUSTER-IP    EXTERNAL-IP    PORT(S)          AGE
kubernetes    ClusterIP     10.12.0.1     <none>         443/TCP          58m
redis         LoadBalancer 10.12.9.51    104.197.66.164 6379:30022/TCP   54s
esam191@cloudshell:~/SOFE4630U-tut3/GKE (silver-course-344506)$ kubectl exec-it redis-56644c686c-vsb5g -- redis-cli
127.0.0.1:6379> auth SOFE4630U
OK
127.0.0.1:6379> set k1 test
OK
127.0.0.1:6379> get k1
"test"
127.0.0.1:6379> set key1 98.26
OK
127.0.0.1:6379> get key1
"98.26"
127.0.0.1:6379> keys *
1) "key1"
2) "k1"
127.0.0.1:6379> keys k?
1) "k1"
127.0.0.1:6379> keys k?y*
1) "key1"
127.0.0.1:6379> set Course "Cloud Computing"
OK
127.0.0.1:6379> get Course
"Cloud Computing"
127.0.0.1:6379> exit
```

```
esam191@cloudshell:~/SOFE4630U-tut3/GKE (silver-course-344506)$ redis-cli -h
104.197.66.164
Unrecognized option or bad number of args for: '-'
esam191@cloudshell:~/SOFE4630U-tut3/GKE (silver-course-344506)$ redis-cli -h 104.197.66.164
104.197.66.164:6379> auth SOFE4630U
OK
104.197.66.164:6379> keys *
1) "key1"
2) "Course"
3) "k1"
104.197.66.164:6379> select 0
OK
104.197.66.164:6379> keys *
1) "key1"
2) "Course"
3) "k1"
104.197.66.164:6379> select 1
OK
104.197.66.164:6379[1]> keys *
(empty array)
104.197.66.164:6379[1]> select 0
OK
104.197.66.164:6379> keys *
1) "key1"
2) "Course"
3) "k1"
104.197.66.164:6379>
```

```
import redis # pip install redis
ip="104.197.66.164"
r = redis.Redis(host=ip, port=6379, db=0,password='SOFE4630U')
v=r.get('key1');
print(v);
r.set('key1','30'.encode('utf-8'));
~
~
```

```
OK
C:\ Command Prompt
D:\CloudComp\SOFE4630U-tut3\python>py redis_access.py
b'98.26'

D:\CloudComp\SOFE4630U-tut3\python>py redis_access.py
b'30'

D:\CloudComp\SOFE4630U-tut3\python>py redis_access.py
b'30'

D:\CloudComp\SOFE4630U-tut3\python>
esaml91@cloudshell:~/SOFE4630U-tut3/python (silver-course-344506)$ redis-cli
-h 104.197.66.164
104.197.66.164:6379> auth SOFE4630U
OK
104.197.66.164:6379> get key1
"98.26"
104.197.66.164:6379> get k1
"test"
104.197.66.164:6379> get key1
"30"
104.197.66.164:6379> 
```

5. Follow the following video to set up sink and source Kafka connectors to the deployed MySQL database.

Video 4:

<https://www.youtube.com/watch?v=B9sgQOQoGHA&feature=youtu.be>

```
Successfully installed avro-python3-1.10.12
mihirkumar@mihirkumar-Precision-3520:~/Documents/cloudComputing/Milestone3/S0FE4630U-tut3/connectors/mysql$ python3 cons_mysql.py
partition:0
key:
value:
{'id': 1, 'name': 'alice', 'email': 'alice@abc.com', 'department': 'eng.', 'modified': 1647584045000}
-----
partition:0
key:
value:
{'id': 2, 'name': 'bob1', 'email': 'bob1@abc.com', 'department': 'sales', 'modified': 1647584045000}
-----
partition:0
key:
value:
{'id': 3, 'name': 'bob2', 'email': 'bob2@abc.com', 'department': 'sales', 'modified': 1647584045000}
-----
partition:0
key:
value:
{'id': 4, 'name': 'bob3', 'email': 'bob3@abc.com', 'department': 'sales', 'modified': 1647584045000}
-----
partition:0
key:
value:
{'id': 5, 'name': 'bob4', 'email': 'bob4@abc.com', 'department': 'sales', 'modified': 1647584045000}
-----
partition:0
key:
value:
{'id': 6, 'name': 'bob5', 'email': 'bob5@abc.com', 'department': 'sales', 'modified': 1647584045000}
-----
partition:0
key:
value:
{'id': 7, 'name': 'bob6', 'email': 'bob6@abc.com', 'department': 'sales', 'modified': 1647584045000}
-----
partition:0
key:
value:
```



```
value:
{'id': 3, 'name': 'bob2', 'email': 'bob2@abc.com', 'department': 'sales', 'modified': 1647590313000}
-----

partition:0
key:
value:
{'id': 4, 'name': 'bob3', 'email': 'bob3@abc.com', 'department': 'sales', 'modified': 1647590313000}
-----

partition:0
key:
value:
{'id': 5, 'name': 'bob4', 'email': 'bob4@abc.com', 'department': 'sales', 'modified': 1647590313000}
-----

partition:0
key:
value:
{'id': 6, 'name': 'bob5', 'email': 'bob5@abc.com', 'department': 'sales', 'modified': 1647590313000}
-----

partition:0
key:
value:
{'id': 7, 'name': 'bob6', 'email': 'bob6@abc.com', 'department': 'sales', 'modified': 1647590313000}
-----

partition:0
key:
value:
{'id': 8, 'name': 'bob7', 'email': 'bob7@abc.com', 'department': 'sales', 'modified': 1647590313000}
-----

partition:0
key:
value:
{'id': 9, 'name': 'bob8', 'email': 'bob8@abc.com', 'department': 'sales', 'modified': 1647590313000}
-----

partition:0
key:
value:
{'id': 10, 'name': 'bob9', 'email': 'bob9@abc.com', 'department': 'sales', 'modified': 1647590313000}
-----
```

Video 5:

https://www.youtube.com/watch?v=yLxfDw9Yh_A&feature=youtu.be

```
CLOUD SHELL
Terminal (galvanized-sled-344502) x + v
All services normal
Open Editor

pmihir2609@cloudshell:~ (galvanized-sled-344502)$ mysql -uuser -pS0FE4638U -h35.193.165.175<<"select * from myDB.TOMySQL;"
mysql: [Warning] Using a password on the command line interface can be insecure.
id      name      email      department  modified
15      user      test@gmail.com  IT          2022-03-18 14:39:37.845
15      esam      esam@gmail.com  IT          2022-03-18 14:51:45.556
15      mike      mike@gmail.com  IT          2022-03-18 14:52:04.632
pmihir2609@cloudshell:~ (galvanized-sled-344502)$
```

Connector name	Status	Category	ID	Plugin name	Tasks	Messages/sec	Bytes/sec	Messages behind	Message
MySQLSinkConnector_0	Running	Sink	lcc-zmm93z	MySQLSink	1	0	--	0	0
MySQLSourceConnector_0	Running	Source	lcc-pggoqm	MySQLSource	1	0	0B/s	0	0

6. Follow the following video to set up a Kafka connector to the deployed Redis database.

Video 6:

<https://www.youtube.com/watch?v=dJWUsAVigR4&feature=youtu.be>

Connectors


[+ Add Connector](#)


Search connectors


Connector name	Status	Category	ID	Plugin name	Tasks	Messages/sec	Bytes/sec	Messages behind	Message
MySQLSinkConnector_0	Running	Sink	lcc-zmm93z	MySQLSink	1	0	--	0	0
MySQLSourceConnector_0	Running	Source	lcc-pggoqm	MySQLSource	1	0	0B/s	0	0
RedisSinkConnector_0	Provisioning	Sink	lcc-6kk18q	RedisSink	0 → 1	0	--	0	0


Connect with popular connectors


[See all connectors](#)


 **Snowflake Sink**
Sink


 **Google Cloud Storage Sink**
Sink


 **Elasticsearch Service Sink**
Sink

 **MongoDB Atlas Source**
Source

 **MongoDB Atlas Sink**
Sink

 **Microsoft SQL Server CDC Source**
Source

 **Postgres CDC Source**
Source

 **Amazon Kinesis Source**
Source

2

[illegible]

9. List some possible applications that can be implemented by using the uploaded dataset.

- In the agricultural sense, it is able to define borders for maps
- Data analysis for companies
- Foot traffic for websites
- It is able to store real time deliveries for products.