Project Milestone

Data Storage Implementation: KV + relational

Name: Mihir Patel
Student Number: 100702168
Date: January 31, 2022

Procedure:

1. **Watch the first three videos for Kafka connectors (focus on the concepts, not the details)**
    - https://www.confluent.io/blog/kafka-connect-tutorial/ (Links to an external site.).

2. **Describe the following:**
    - Sink and Source connectors.
    - The applications/advantages of using Kafka Connectors with data storage.
    - How do Kafka connectors maintain availability?
    - List the popular Kafka converters for values and the properties/advantages of each.

3. **Search the internet to answer the following question:**
    - What's a Key-Value (KV) database?
    - What are KV databases' advantages and disadvantages?
    - List some popular KV databases.

4. **Follow the following videos to deploy and use Redis and MySQL databases using GKE.**

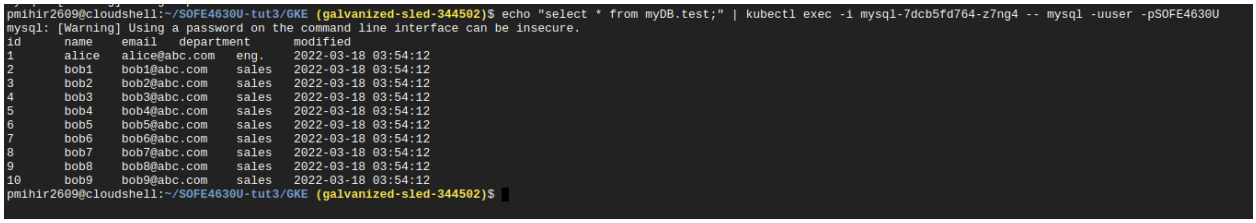

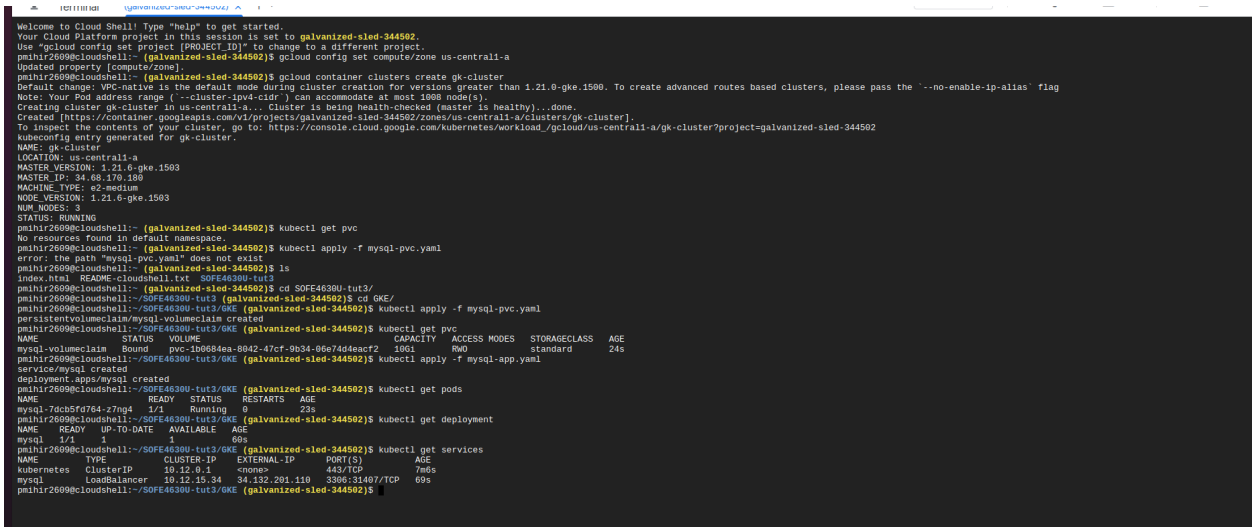

```
pmihir2609@cloudshell:~/SOFE4630U-tut3/GKE (galvanized-sled-344502)$ kubectl delete -f mysql.app.yaml
error: the path "mysql.app.yaml" does not exist
pmihir2609@cloudshell:~/SOFE4630U-tut3/GKE (galvanized-sled-344502)$ kubectl delete -f mysql-app.yaml
service "mysql" deleted
deployment.apps "mysql" deleted
pmihir2609@cloudshell:~/SOFE4630U-tut3/GKE (galvanized-sled-344502)$ kubectl apply -f mysql-app.yaml
service/mysql created
deployment.apps/mysql created
pmihir2609@cloudshell:~/SOFE4630U-tut3/GKE (galvanized-sled-344502)$ kubectl get services
NAME         TYPE           CLUSTER-IP    EXTERNAL-IP   PORT(S)          AGE
kubernetes   ClusterIP      10.12.0.1     <none>        443/TCP          52m
mysql        LoadBalancer   10.12.0.202   <pending>     3306:32255/TCP   16s
pmihir2609@cloudshell:~/SOFE4630U-tut3/GKE (galvanized-sled-344502)$ kubectl get services
NAME         TYPE           CLUSTER-IP    EXTERNAL-IP   PORT(S)          AGE
kubernetes   ClusterIP      10.12.0.1     <none>        443/TCP          52m
mysql        LoadBalancer   10.12.0.202   <pending>     3306:32255/TCP   29s
pmihir2609@cloudshell:~/SOFE4630U-tut3/GKE (galvanized-sled-344502)$ kubectl get services
NAME         TYPE           CLUSTER-IP    EXTERNAL-IP      PORT(S)          AGE
kubernetes   ClusterIP      10.12.0.1     <none>           443/TCP          52m
mysql        LoadBalancer   10.12.0.202   35.193.165.175   3306:32255/TCP   55s
pmihir2609@cloudshell:~/SOFE4630U-tut3/GKE (galvanized-sled-344502)$ mysql -uuser -pSOFE4630U -h35.193.165.175 <<<" select * form  myDB.test;"
mysql: [Warning] Using a password on the command line interface can be insecure.
ERROR 1064 (42000) at line 1: You have an error in your SQL syntax; check the manual that corresponds to your MySQL server version for the right syntax to use near 'form  myDB.test' at line 1
pmihir2609@cloudshell:~/SOFE4630U-tut3/GKE (galvanized-sled-344502)$ mysql -uuser -pSOFE4630U -h35.193.165.175 <<<"select * from  myDB.test;"
mysql: [Warning] Using a password on the command line interface can be insecure.
id      name    email    department      modified
1       alice   alice@abc.com   eng.    2022-03-18 04:04:31
2       bob1    bob1@abc.com    sales   2022-03-18 04:04:31
3       bob2    bob2@abc.com    sales   2022-03-18 04:04:31
4       bob3    bob3@abc.com    sales   2022-03-18 04:04:31
5       bob4    bob4@abc.com    sales   2022-03-18 04:04:31
6       bob5    bob5@abc.com    sales   2022-03-18 04:04:31
7       bob6    bob6@abc.com    sales   2022-03-18 04:04:31
8       bob7    bob7@abc.com    sales   2022-03-18 04:04:31
9       bob8    bob8@abc.com    sales   2022-03-18 04:04:31
10      bob9    bob9@abc.com    sales   2022-03-18 04:04:31
pmihir2609@cloudshell:~/SOFE4630U-tut3/GKE (galvanized-sled-344502)$ kubectl delete -f mysql-app.yaml
service "mysql" deleted
deployment.apps "mysql" deleted
pmihir2609@cloudshell:~/SOFE4630U-tut3/GKE (galvanized-sled-344502)$
```

5. **Follow the following video to set up sink and source Kafka connectors to the deployed MySQL database**

```
mihirkumar@mihirkumar-Precision-3520:~/Documents/cloudComputing/Milestone3/SOFE4630U-tut3/connectors/mysql$ python3 cons_mysql.py
partition:0
key:
value:
{'id': 1, 'name': 'alice', 'email': 'alice@abc.com', 'department': 'eng.', 'modified': 1647584045000}

--------

partition:0
key:
value:
{'id': 2, 'name': 'bob1', 'email': 'bob1@abc.com', 'department': 'sales', 'modified': 1647584045000}

--------

partition:0
key:
value:
{'id': 3, 'name': 'bob2', 'email': 'bob2@abc.com', 'department': 'sales', 'modified': 1647584045000}

--------

partition:0
key:
value:
{'id': 4, 'name': 'bob3', 'email': 'bob3@abc.com', 'department': 'sales', 'modified': 1647584045000}

--------

partition:0
key:
value:
{'id': 5, 'name': 'bob4', 'email': 'bob4@abc.com', 'department': 'sales', 'modified': 1647584045000}

--------

partition:0
key:
value:
{'id': 6, 'name': 'bob5', 'email': 'bob5@abc.com', 'department': 'sales', 'modified': 1647584045000}

--------

partition:0
key:
value:
{'id': 7, 'name': 'bob6', 'email': 'bob6@abc.com', 'department': 'sales', 'modified': 1647584045000}

--------

partition:0
key:
value:
```

```
value:
{'id': 3, 'name': 'bob2', 'email': 'bob2@abc.com', 'department': 'sales', 'modified': 1647590313000}

--------

partition:0
key:
value:
{'id': 4, 'name': 'bob3', 'email': 'bob3@abc.com', 'department': 'sales', 'modified': 1647590313000}

--------

partition:0
key:
value:
{'id': 5, 'name': 'bob4', 'email': 'bob4@abc.com', 'department': 'sales', 'modified': 1647590313000}

--------

partition:0
key:
value:
{'id': 6, 'name': 'bob5', 'email': 'bob5@abc.com', 'department': 'sales', 'modified': 1647590313000}

--------

partition:0
key:
value:
{'id': 7, 'name': 'bob6', 'email': 'bob6@abc.com', 'department': 'sales', 'modified': 1647590313000}

--------

partition:0
key:
value:
{'id': 8, 'name': 'bob7', 'email': 'bob7@abc.com', 'department': 'sales', 'modified': 1647590313000}

--------

partition:0
key:
value:
{'id': 9, 'name': 'bob8', 'email': 'bob8@abc.com', 'department': 'sales', 'modified': 1647590313000}

--------

partition:0
key:
value:
{'id': 10, 'name': 'bob9', 'email': 'bob9@abc.com', 'department': 'sales', 'modified': 1647590313000}

--------
```

6. **Follow the following video to set up a Kafka connector to the deployed Redis database.**

Connector name	Status	Category	ID	Plugin name	Tasks	Messages/sec	Bytes/sec	Messages behind	Messag
MySqlSinkConnector_0	Running	Sink	lcc-zmm93z	MySqlSink	1	0	--	0	0
MySqlSourceConnector_0	Running	Source	lcc-pggoqm	MySqlSource	1	0	0B/s	0	0

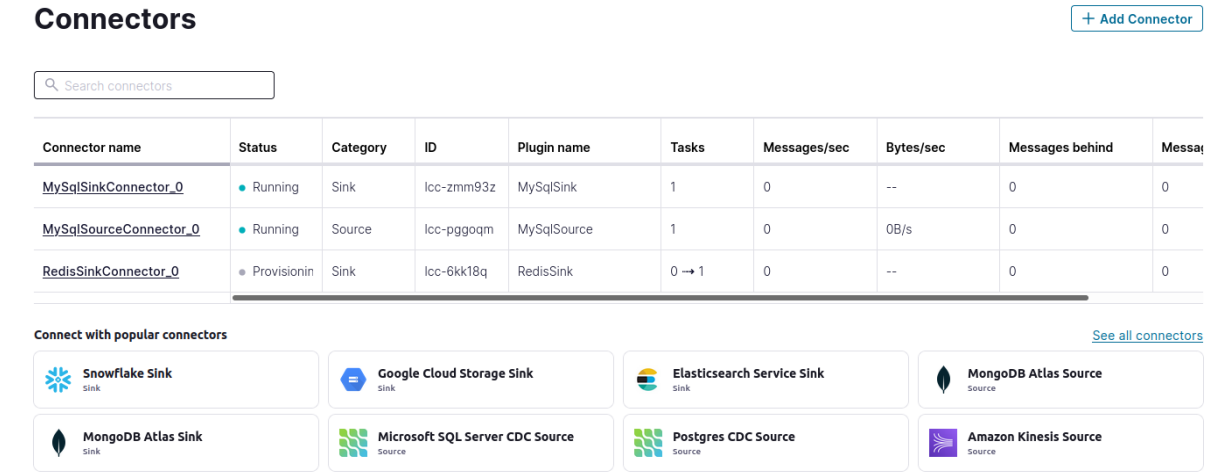

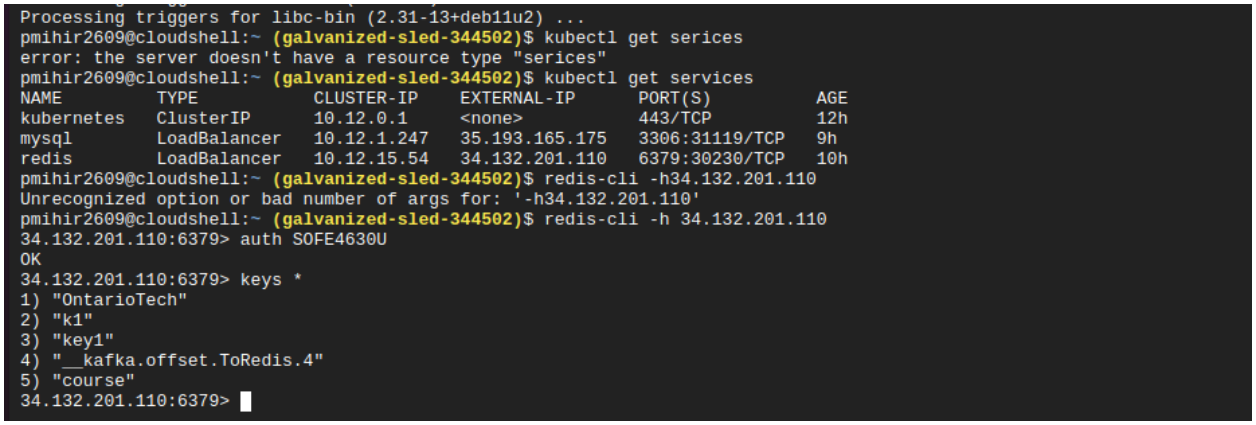

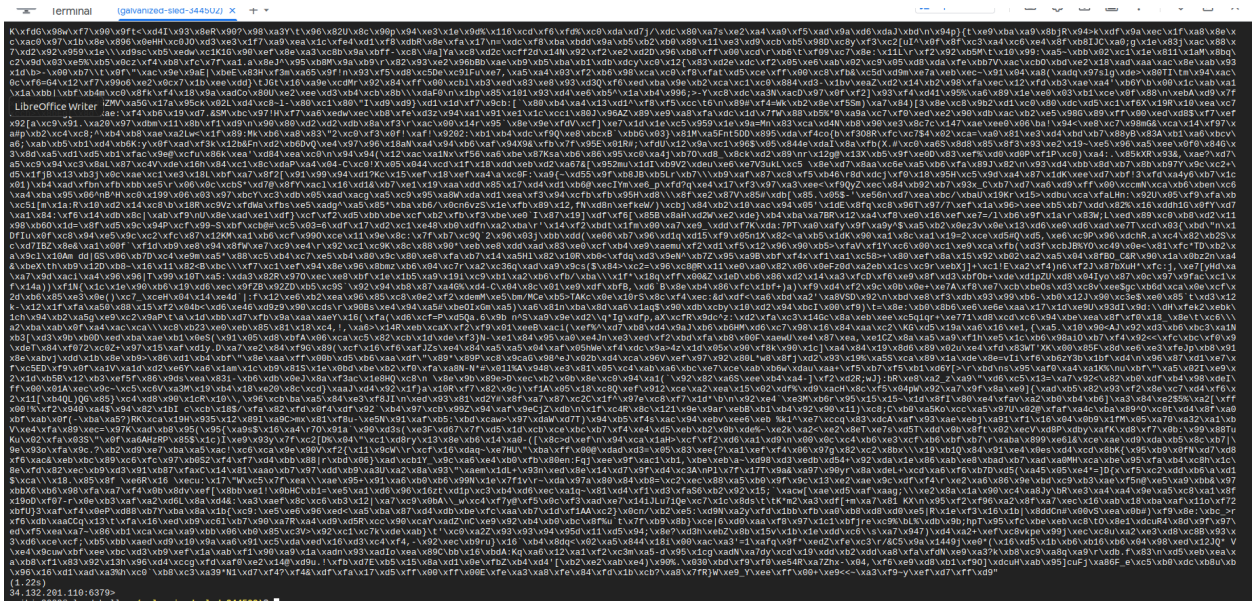

7. **Now, you will store a dataset into cloud storage. The dataset has to be sent into Kafka topics and connectors have to be configured to automatically store the dataset into the dataset storage. The producer that will send the dataset to Kafka topics should run on your local machine as it will simulate real sensors while Kafka, connectors, and data storage should be on the cloud. Use MySQL for the CSV files and Redis for images. Feel free to update the Yaml files from the given repository to fit your dataset**

8. **Use the sensors, images, ground truth Pose in the lastest session in http://robots.engin.umich.edu/nclt/ (Links to an external site.) as your dataset.**

9. **List some possible applications that can be implemented by using the uploaded dataset.**